

SIT 788: Assignment 5.2

Computer Vision and Custom Vision

Part 1: Real time face detection application

This part has walkthrough for creating a real time face detection application. The resource used (Azure Face detection API [1]) is same from the 5.1 Task.

Step 1: Creating face detection resource

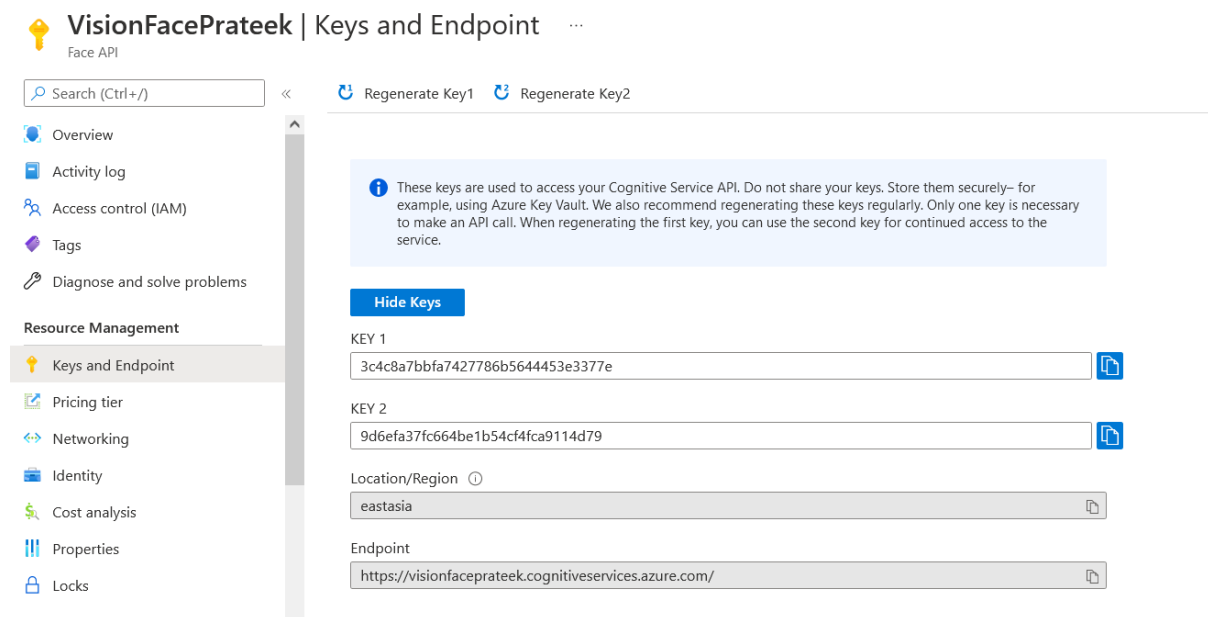


Fig 1.1: Resource used for face detection.

Step 2: Import libraries and connect to the resource

Here we import the libraries and connect to the face detection resource using the generated subscription key and endpoint.

Part 1: Real-time face detection

```
In [1]: from msrest.authentication import CognitiveServicesCredentials
from azure.cognitiveservices.vision.face import FaceClient

import cv2
from IPython.display import Image
from matplotlib import pyplot as plt
from time import time

In [2]: attributes_list = ['age', 'emotion', 'gender']
fps = 1
fps = fps * 1000

subscription_key = "3c4c8a7bbfa7427786b5644453e3377e"
endpoint = "https://visionfaceprateek.cognitiveservices.azure.com/"
cv_face_client = FaceClient(
    endpoint, CognitiveServicesCredentials(subscription_key))
print("Connected to resource", cv_face_client)

Connected to resource <azure.cognitiveservices.vision.face._face_client.FaceClient object at 0x000001C4FFD48040>
```

Fig 2.1: Import libs and connect to the face detection resource

Step 3: Define function to process passed frame and process results

Using the resource, a function `process_frame` is defined which sends the frame to the azure face detection API. A predefined attribute list is also used to generate additional information (age, emotion, and gender).

```
In [3]: def process_frame(frame2api):
        cv2.imwrite('tmp_image.jpg', frame2api)
        tmp_img = open("tmp_image.jpg", 'rb')
        detected_faces = cv_face_client.face.detect_with_stream(
            tmp_img, return_face_attributes=attributes_list)
        tmp_img.close()
        return detected_faces
```

Fig 3.1: Function to detect faces

Another function to process the results and overlay annotations on detection frame is defined as *process results*.

```
In [4]: font = cv2.FONT_HERSHEY_SIMPLEX
        fontScale = 1
        color = (0, 0, 255)
        thickness = 2

        def process_result(frame2api, detected_faces):
            if len(detected_faces) > 0:
                print("Number of faces detected:", len(detected_faces))
                for faces_itr in range(len(detected_faces)):
                    org = (detected_faces[faces_itr].face_rectangle.left,
                           detected_faces[faces_itr].face_rectangle.top)
                    gender = detected_faces[faces_itr].face_attributes.gender.value
                    age = detected_faces[faces_itr].face_attributes.age
                    emotions_dict = detected_faces[faces_itr].face_attributes.emotion.as_dict(
                    )
                    emotion_name = max(
                        zip(emotions_dict.values(), emotions_dict.keys()))[1]

                    face_rect = [detected_faces[faces_itr].face_rectangle.left,
                                 detected_faces[faces_itr].face_rectangle.top,
                                 detected_faces[faces_itr].face_rectangle.width,
                                 detected_faces[faces_itr].face_rectangle.height]

                    frame2api = cv2.rectangle(frame2api, face_rect, (0, 0, 255), 5)
                    frame2api = cv2.putText(frame2api, str(
                        gender), org, font, fontScale, color, thickness, cv2.LINE_AA)
                    frame2api = cv2.putText(frame2api, str(
                        age), (org[0]+100, org[1]), font, fontScale, color, thickness, cv2.LINE_AA)
                    frame2api = cv2.putText(frame2api, str(
                        emotion_name), (org[0] + 200, org[1]), font, fontScale, color, thickness, cv2.LINE_AA)
            else:
                print("No Face Detected")
            return len(detected_faces)
```

Fig 3.2: process_result() overlays age, gender and emotion information along with face bbox on the processed framed.

Step 4: Running live camera using OpenCV

OpenCV is used to initialize a live video stream using the webcam. As the purpose of the task is also to sample the frames such that 1 frame per second is passed, time difference between two successive processed frame is calculated. This ensures that there is at least 1 second difference between two consecutive processed frame, rest frames are skipped. This also ensures that the API is called efficiently and, in a cost-effective way.

```
In [5]: vid = cv2.VideoCapture(0)
start_time = time() * 1000
init_frame = True

while (True):
    _, frame = vid.read()
    if init_frame:
        frame2api = frame.copy()
        detected_faces = process_frame(frame2api)
        process_result(frame2api, detected_faces)
        init_frame = False

    end_time = time() * 1000
    time_diff = end_time - start_time
    if (int(time_diff) >= fps):
        frame2api = frame.copy()
        detected_faces = process_frame(frame2api)
        process_result(frame2api, detected_faces)
        start_time = time() * 1000

    frame = cv2.putText(frame, "Real Time", (10,10), font,
                        fontScale, color, thickness, cv2.LINE_AA)
    frame2api = cv2.putText(frame2api, "Processed", (50,50),
                            font, fontScale, color, thickness, cv2.LINE_AA)
    frame_merged = cv2.hconcat([frame, frame2api])
    cv2.imshow("Display", frame_merged)

    if cv2.waitKey(1) & 0xFF == ord('q'):
        break
vid.release()
cv2.destroyAllWindows()

Number of faces detected: 1
Number of faces detected: 1
Number of faces detected: 1
Number of faces detected: 1
Number of faces detected: 1
Number of faces detected: 1
Number of faces detected: 1
Number of faces detected: 1
```

Fig 4.1: Using OpenCV to run video stream via webcam. 1 FPS is set as the default parameter.




Fig 4.2: Output showing real time frame [Left] vs Last processed frame with face bbox and facial attributes overlayed [Right]




Fig 4.3: Output showing real time frame [Left] vs Last processed frame with face bbox and facial attributes overlayed [Right]

Part 2: Custom Vision

Custom Vision [3] is offered by Azure for the scenario when we need to use our own data and train a model from scratch. Currently, object detection and object classification are supported in this service. In this experiment, we use Dogs and Cats dataset [2] to train a classifier using Azure Custom Vision. This model is then deployed and tested using Azure Custom Vision.

Step 1: Creating a custom vision resource

Once the test resource is created via azure portal, we can see two resources: one for training and other for prediction.

CustomVisionPrateek | Keys and Endpoint ...

Custom vision

Search (Ctrl+ /) << Regenerate Key1 Regenerate Key2

Overview
Activity log
Access control (IAM)
Tags
Diagnose and solve problems

Resource Management

Keys and Endpoint
Encryption
Pricing tier
Networking
Identity
Cost analysis
Properties
Locks

Monitoring

These keys are used to access your Cognitive Service API. Do not share your keys. Store them securely– for example, using Azure Key Vault. We also recommend regenerating these keys regularly. Only one key is necessary to make an API call. When regenerating the first key, you can use the second key for continued access to the service.

Hide Keys

KEY 1
ff801d0322034b64b8091684cf72db55

KEY 2
8322e59a051d47169b1fb583ca4709a7

Location/Region ⓘ
centralindia

Endpoint
https://customvisionprateek.cognitiveservices.azure.com/

(a)

CustomVisionPrateek-Prediction | Keys and Endpoint ...

Custom vision

Search (Ctrl+ /) << Regenerate Key1 Regenerate Key2

Overview
Activity log
Access control (IAM)
Tags
Diagnose and solve problems

Resource Management

Keys and Endpoint
Encryption
Pricing tier
Networking
Identity
Cost analysis
Properties

These keys are used to access your Cognitive Service API. Do not share your keys. Store them securely– for example, using Azure Key Vault. We also recommend regenerating these keys regularly. Only one key is necessary to make an API call. When regenerating the first key, you can use the second key for continued access to the service.

Hide Keys

KEY 1
e6c1b0b739da430fa32d61e59ed8ef21

KEY 2
17bebdacd2a548ef84940fd3992e1211

Location/Region ⓘ
centralindia

Endpoint
https://customvisionprateek-prediction.cognitiveservices.azure.com/

(b)

Fig 1.1: (a) Keys and endpoint for training resource (b) Keys and endpoint for prediction resource.

Step 2: Installing and importing dependencies

After resource creation, python dependencies for custom vision are installed.

```
In [6]: !pip install azure-cognitiveservices-vision-customvision

Requirement already satisfied: azure-cognitiveservices-vision-customvision in c:\users\singh\anaconda3\lib\site-packages (3.1.0)
Requirement already satisfied: azure-common~=1.1 in c:\users\singh\anaconda3\lib\site-packages (from azure-cognitiveservices-vision-customvision) (1.1.28)
Requirement already satisfied: msrest>=0.5.0 in c:\users\singh\anaconda3\lib\site-packages (from azure-cognitiveservices-vision-customvision) (0.6.21)
Requirement already satisfied: isodate>=0.6.0 in c:\users\singh\anaconda3\lib\site-packages (from msrest>=0.5.0->azure-cognitiveservices-vision-customvision) (0.6.1)
Requirement already satisfied: requests-oauthlib>=0.5.0 in c:\users\singh\anaconda3\lib\site-packages (from msrest>=0.5.0->azure-cognitiveservices-vision-customvision) (1.3.1)
Requirement already satisfied: requests~=2.16 in c:\users\singh\anaconda3\lib\site-packages (from msrest>=0.5.0->azure-cognitiveservices-vision-customvision) (2.25.1)
Requirement already satisfied: certifi>=2017.4.17 in c:\users\singh\anaconda3\lib\site-packages (from msrest>=0.5.0->azure-cognitiveservices-vision-customvision) (2020.12.5)
Requirement already satisfied: six in c:\users\singh\anaconda3\lib\site-packages (from isodate>=0.6.0->msrest>=0.5.0->azure-cognitiveservices-vision-customvision) (1.15.0)
Requirement already satisfied: chardet<5,>=3.0.2 in c:\users\singh\anaconda3\lib\site-packages (from requests~=2.16->msrest>=0.5.0->azure-cognitiveservices-vision-customvision) (4.0.0)
Requirement already satisfied: idna<3,>=2.5 in c:\users\singh\anaconda3\lib\site-packages (from requests~=2.16->msrest>=0.5.0->azure-cognitiveservices-vision-customvision) (2.10)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in c:\users\singh\anaconda3\lib\site-packages (from requests~=2.16->msrest>=0.5.0->azure-cognitiveservices-vision-customvision) (1.26.4)
Requirement already satisfied: oauthlib>=3.0.0 in c:\users\singh\anaconda3\lib\site-packages (from requests-oauthlib>=0.5.0->msrest>=0.5.0->azure-cognitiveservices-vision-customvision) (3.2.0)

In [153]: from azure.cognitiveservices.vision.customvision.training import CustomVisionTrainingClient
from azure.cognitiveservices.vision.customvision.prediction import CustomVisionPredictionClient
from azure.cognitiveservices.vision.customvision.training.models import ImageFileCreateBatch, ImageFileCreateEntry, Region
from msrest.authentication import ApiKeyCredentials
import os, time, uuid
```

Fig 2.1: Installing custom vision and importing libs

Step 3: Connecting to the resource

The resource created in step 1 is connected via subscription key and endpoints. Two clients, one for training and one for prediction are created.

Note: The subscription keys and endpoints for prediction and training are different, hence they are connected separately.

```
In [154]: endpoint = "https://customvisionprateek.cognitiveservices.azure.com/"
training_key = "ff801d0322034b64b8091684cf72db55"
prediction_key = "e6c1b0b739da430fa32d61e59ed8ef21"
prediction_resource_id = "/subscriptions/d5d0e7cd-5900-4c80-b24c2a8416d0/resourceGroups/PRATEEK-SINGH/providers/Microsoft.CognitiveServices/VisionCustomVisionAccounts/PRATEEK-SINGH"

In [155]: credentials = ApiKeyCredentials(in_headers={"Training-key": training_key})
trainer = CustomVisionTrainingClient(endpoint, credentials)
prediction_credentials = ApiKeyCredentials(in_headers={"Prediction-key": prediction_key})
predictor = CustomVisionPredictionClient(endpoint, prediction_credentials)

In [156]: publish_iteration_name = "classifyDogCat-v3"

credentials = ApiKeyCredentials(in_headers={"Training-key": training_key})
trainer = CustomVisionTrainingClient(endpoint, credentials)

# Create a new project
print ("Creating project...")
project_name = uuid.uuid4()
project = trainer.create_project(project_name)

Creating project...
```

Fig 3.1: Clients for training and prediction. Project created using generated credentials.

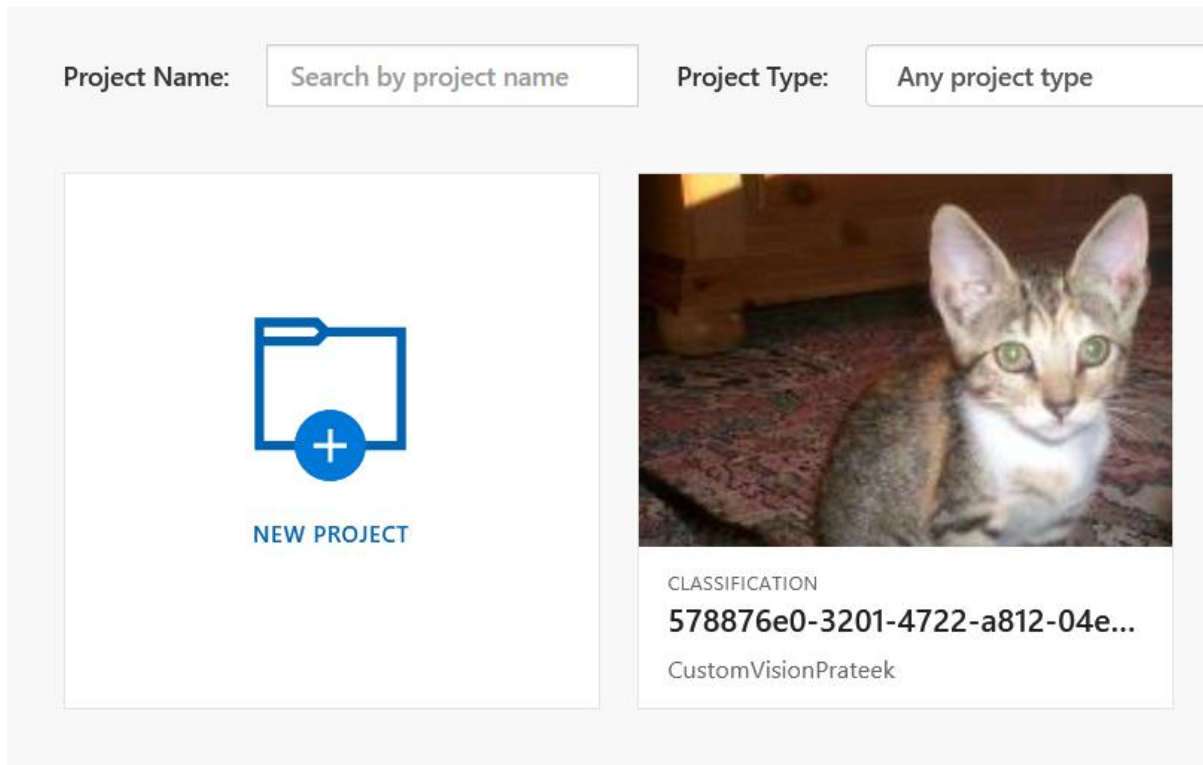


Fig 3.2: Once the project is created using the SDK, it is reflected in the custom vision ai portal.

Step 4: Uploading data and training the classifier

After the project is created, the images are uploaded using the SDK [the SDK only allows 64 images to be uploaded, once the limit is reached then the images can be uploaded again but not in the same go; no such limitation is on the custom vision webpage].

Once the images are uploaded the training is done and model is published on the custom vision portal.

Train using the downloaded dataset

```
In [159]: base_image_location = os.path.dirname("C:\\Users\\singh\\Documents\\MastersAppliedAI\\Deakin\\\"
          \"SIT788_Engineering_AI_Solutions\\Assignments\\SIT788_5_2_Data\\dataset\\training_set\\\"

print("Adding images...")

image_list = []
step = 32
itr = 32
#for itr in range(0, 100, step):
for image_num in range(itr, itr+step):
    file_name = "dog.{0}.jpg".format(image_num+1)
    with open(os.path.join(base_image_location, "dogs", file_name), "rb") as image_contents:
        image_list.append(ImageFileCreateEntry(name=file_name, contents=image_contents.read(), tag_ids=[dog_tag.id]))

for image_num in range(itr, itr+step):
    file_name = "cat.{0}.jpg".format(image_num+1)
    with open(os.path.join(base_image_location, "cats", file_name), "rb") as image_contents:
        image_list.append(ImageFileCreateEntry(name=file_name, contents=image_contents.read(), tag_ids=[cat_tag.id]))

upload_result = trainer.create_images_from_files(project.id, ImageFileCreateBatch(images=image_list))
if not upload_result.is_batch_successful:
    print("Image batch upload failed.")
    for image in upload_result.images:
        print("Image status: ", image.status)
    exit(-1)
else:
    print("Adding images successful")
```

Adding images...
Adding images successful

Fig 4.1: Uploading data

Step 5: Testing using test data

As the data is two class, it can be treated as a binary classification. The class with higher probability is the predicted class *[The sum of probabilities is always 100, so the dominant class will have probability greater than 50%]*. The testing set from [2] has 1000 samples for dogs and cats each. All these images are sent to the prediction resource and evaluation metrics is calculated for binary classification.

```
In [172]: import glob
all_dogs_images = glob.glob(test_image_location + "/dogs/*")
all_cats_images = glob.glob(test_image_location + "/cats/*")

In [186]: from tqdm import tqdm
TP = 0
FN = 0
TN = 0
FP = 0

for image_dog in tqdm(all_dogs_images):
    with open(os.path.join(test_image_location, image_dog), "rb") as image_contents:
        results = predictor.classify_image(
            project.id, publish_iteration_name, image_contents.read())

        for prediction in results.predictions:
            if prediction.probability > 0.5:
                prediction_name = prediction.tag_name
            if (prediction_name == "Dog"):
                TP += 1
            else:
                FN += 1

for image_cat in tqdm(all_cats_images):
    with open(os.path.join(test_image_location, image_cat), "rb") as image_contents:
        results = predictor.classify_image(
            project.id, publish_iteration_name, image_contents.read())

        for prediction in results.predictions:
            if prediction.probability > 0.5:
                prediction_name = prediction.tag_name
            if (prediction_name == "Cat"):
                TN += 1
            else:
                FP += 1

100%|████████████████████████████████████████████████████████████████████████████████| 1000/1000 [07:12<00:00, 2.31it/s]
100%|████████████████████████████████████████████████████████████████████████████████| 1000/1000 [07:06<00:00, 2.34it/s]
```

Fig 5.1: Testing cats and dogs' folder with the trained classifier.

```
In [178]: Recall = TP / (TP + FN)
Precision = TP / (TP + FP)
F_Measure = (2 * Precision * Recall) / (Precision + Recall)
Accuracy = (TP + TN) / (TP + TN + FP + FN)

In [185]: total = TP + TN + FP + FN
print ("Total classified samples: ", total)
print ("Evaluation metrics: Recall: ", Recall, "Precision: ", Precision,
        "F-Measure: ", F_Measure, "Accuracy: ", Accuracy)

Total classified samples: 2000
Evaluation metrics: Recall: 0.981 Precision: 0.9969512195121951 F-Measure: 0.9889112903225806 Accuracy: 0.989
```

Fig 5.2: Evaluation metrics for the trained classifier.

The classifier is trained well considering only 64 images per class were used to train the classifier. The evaluation metrics for the classifier are:

- Precision: 99.6%
- Recall: 98.1%
- f-Measure: 98.9%
- Accuracy: 98.9%

References:

- [1] <https://azure.microsoft.com/en-us/services/cognitive-services/face/>
- [2] <https://www.kaggle.com/datasets/chetankv/dogs-cats-images>
- [3] <https://azure.microsoft.com/en-us/services/cognitive-services/custom-vision-service/>